

ARBITER — What We Built, Why, and What Comes Next

Adaptive Multi-Agent System for Iterative Task Generation, Validation and Refinement

Gautam Lasgotra (23BCS032) · Chirag Attri (23BCS024) · Aniket Kundal (23BCS015)

Guide: Dr. Sonika Gupta, SoCSE · SMVDU Katra · 3 October 2026

Repository: github.com/igautamlasgotra/arbitr · Live: arbitr-gules-two.vercel.app

1. The research objective

1.1 The question

Under a matched token budget, does adaptive multi-agent refinement actually improve correctness over a single agent — and does the answer depend on how strong the available automatic verifier is?

1.2 Why this question and not "can agents build software"

Multi-agent refinement is not new. Neither is adaptive orchestration — there is a published survey of it. **This project claims novelty in neither**, and says so in the synopsis. Claiming otherwise would not survive a viva.

What is genuinely unresolved is narrower and more interesting. Two bodies of work disagree:

Position	Evidence
Multi-agent refinement helps	Self-Refine, Reflexion, AgentCoder, ChatDev, MetaGPT
It mostly does not, once you pay for it	Tran & Kiela (2025): at equal <i>thinking-token</i> budget, one agent matches or beats multi-agent systems on reasoning

Both can be true at once — if the benefit depends on whether a cheap, trustworthy **verifier** exists for the task. Huang et al. (ICLR 2024) supply the mechanism: without external feedback, a model cannot reliably correct itself.

Here is the gap we occupy. Adaptive orchestration's entire justification is "*spend agents only where they are needed*" — a claim about **compute efficiency**. It has therefore never been tested by the one protocol that measures compute efficiency directly: **budget matching**. We run that test, across three task families chosen to differ in verifier strength.

1.3 Why this is safe to attempt on a deadline

A **negative result** is a valid result. "We built it, held the budget equal, and adaptivity did not pay off on family X" is a finding, not a failure. No other framing available to us has that property.

1.4 The questions, stated precisely

- **RQ1** Under a matched token budget, does adaptive orchestration beat (a) a budget-matched single agent and (b) fixed multi-agent pipelines?
- **RQ2** Does any benefit track **verifier strength** (executable tests → execution match → answer only)?
- **RQ3** Does combining a tool validator with an LLM critic reduce **false accepts** compared with LLM-only review?
- **RQ4** How many refinement iterations stay useful before returns vanish?

2. What the system is — and what it is not

ARBITER takes a task in plain English, decides which agent roles to use, produces a solution, **validates it by running it**, feeds the failures back, and revises — until it passes or a budget stops it. Every step is recorded.

2.1 What it deliberately does not do

It does not build applications, games, websites or mobile apps. This is not a missing feature; it is a direct consequence of the research question.

The loop works because the validator returns an **objective score**. Ask for a Flutter app and nothing can produce that score: there is no single return value to assert on, and "is this calculator good?" has no automatic answer. Without a verifier the loop degrades into one model asking another model for an opinion — **which is precisely the thing Huang et al. showed does not work, and precisely what this project is gathering evidence about.**

So the system **declines** tasks it cannot check, in plain words, rather than appearing to validate them. Knowing the boundary of what it can verify is a property of the design, not an apology for it.


3. The logic — why each part exists

3.1 An agent is a function, not a framework

```
def generator(task, router, state, feedback=None):
    prompt = build_prompt(task, feedback)
    resp = router.call(prompt, schema=SOLUTION_SCHEMA, state=state)
    return Solution.model_validate_json(resp.text)
```

A role prompt, one model call, a typed parse. Agents differ **only in their system prompt**: `repair` is told not to trust the previous approach, `critic` is told to review sceptically. No LangGraph, no AutoGen — the orchestration logic is what the project studies, so burying it inside a framework would hide the contribution. (AutoGen is also in maintenance mode.)

3.2 The orchestrator is a bounded loop



It can never run unbounded. It stops on acceptance, 5 iterations, 20 model calls, 60,000 tokens, 180 seconds, two iterations without improvement, or a repeated identical output. **Every stop reason is logged as an enum.** The last condition matters: step repetition is the single most common documented multi-agent failure mode — 17.1% of failures in the MAST taxonomy (NeurIPS 2025).

3.3 Validation contains no LLM

For code, `python_exec` writes the program and its tests to a temp directory, runs them in a **separate process with no inherited environment**, and reports which named tests failed. That subprocess is the external feedback the entire loop depends on.

3.4 The score is graded, not binary

`0.80` means four of five tests passed. Without partial credit the system could not tell "getting closer" from "stuck", and the no-improvement stop condition would be meaningless. It also makes defects visible — see §5.3.

3.5 The test-designer agent

Benchmark tasks ship with tests. A task typed by a panel member does not. The test designer derives executable tests from the description **and fixes the function signature**, which is then handed to the generator — otherwise the tester

invents `two_sum` while the generator writes `twoSum` and everything fails for the wrong reason.

Generated tests are parsed and checked before use: they must define `check_*` functions, must call the declared entry point, must be valid Python, and may import only from a pure-standard-library allowlist. `random` and `time` are excluded deliberately — a non-deterministic test destroys the refinement signal.

Stated limitation: model-written tests are not ground truth. They are used for live unseen tasks only. All benchmark measurement uses the dataset's own tests, and the two are never mixed in results.

3.6 The budget governor — the methodological heart

Every model call is metered: prompt tokens, completion tokens, calls, wall time. This is not instrumentation added at the end. It is what makes a budget-matched comparison possible at all, and therefore what the research contribution rests on.

3.7 The five conditions

ID	Condition	Roles on screen
A	Single agent, one shot	generator
A+	Single agent given the same token budget as D, spent on repeated sampling with majority voting	generator ×N
B	Generator ↔ validator refinement loop	generator
C	Fixed pipeline	generator → critic → repair
D	ARBITER adaptive — planner picks roles and depth per task	varies

A+ is what makes the project credible. Most student multi-agent projects compare against a single call, declare victory, and ignore that the multi-agent system spent five times the tokens. A+ removes that confound. It is deliberately **denied the validator** — if it could see test results it would be a refinement loop, not a single agent, and the comparison would be rigged.

All five share one engine and differ only by configuration, so a difference can never come from implementation divergence.

4. How we built it

4.1 Stack

Layer	Choice	Reason
Language	Python 3.10	Local install; ecosystem
API	FastAPI + Server-Sent Events	One-directional stream, self-reconnecting, far less code than WebSockets
Types	Pydantic v2	Forces structured model output into typed objects
Orchestration	Custom, ~400 lines	The object of study must stay readable
Model	Google Gemini Flash-Lite via raw REST	Exact control over the request body, real token counts
Sandbox	Subprocess, stripped env, timeout	Generated code must not reach API keys
Storage	JSONL traces + on-disk response cache	Every metric is computed from traces; no second pipeline
Hosting	Vercel (fluid compute)	Fluid compute is what lets one request stream for a whole run

4.2 Layout

arbiter/		
├ core/	schemas.py, state.py	← frozen contracts + budget governor
├ llm/	base, cache, gemini, router, mock	
├ agents/	roles.py, test_designer.py	
├ validators/	python_exec.py, answer_match.py	
├ sandbox/	runner.py	← isolated subprocess
├ orchestrator/	loop.py, planner.py, baselines.py	
├ web/	app.py, templates/	← FastAPI + SSE
└ bench/	runner.py, metrics.py, datasets/	

4.3 Decisions worth defending

Decision	Alternative	Why
Role selection from a fixed library	Inventing agents at runtime	Runtime invention is unbounded and unreproducible — a demo, not an experiment
Tools validate, LLM critic only advises	LLM-only review	Self-correction without external feedback is unreliable; the critic's verdict is logged separately to measure false accepts
Disk response cache from day one	Add caching later	Re-runs cost nothing and become byte-reproducible; without it the free-tier quota makes the experiment infeasible
Graded score	Binary pass/fail	Partial credit is what detects improvement and stagnation
Raw REST, no Gemini SDK	Official SDK	The cache key depends on the exact request body; also gives real <code>usageMetadata</code> counts
Token-gated live runs on the public host	Open endpoint	A public URL that executes model-written Python is a remote shell

4.4 Quota strategy

A full sweep is several thousand calls against a free tier. Four controls: the **response cache** (52% hit rate on the first sweep, so re-runs are largely free), **round-robin across keys** with per-key daily-quota retirement, a **resumable runner** that checkpoints after every task-run and continues exactly where a quota error stopped it, and a **reserved demo key** the experiment runner never touches, so demo-day quota is always fresh.

5. What is working today

5.1 Verified live behaviour

Run on the deployed instance, task typed at runtime, condition C:

test_designer	entry=solve, 5 tests	
generator	solution produced	
python_exec	score=0.80	← FAILED: 4 of 5 tests passed
llm_critic	passed=True	← the LLM said it was fine
decide	refine	
repair	new solution (a different agent)	
python_exec	score=1.00	← accepted

ARBITER.

Adaptive Role-Based Iterative Task Execution & Refinement

live mode ready

Task — type anything, this is not a scripted demo

Evaluate a string arithmetic expression containing non-negative integers and the operators + - * / with correct operator precedence and no parentheses. Division truncates toward zero. Do not use eval.

Family

Condition

code — validated by running tests

B — generator ↔ validator loop

Run

Offline demo

TASK condition B

Evaluate a string arithmetic expression containing non-negative integers and the operators + - * / with correct operator precedence and no parentheses. Division truncates toward zero. Do not use eval.

TEST_DESIGNER entry=evaluate_expression, 5 tests 1.49s

PLAN baseline B 1.49s

GENERATOR iter 1 1357 chars 3.46s

PYTHON_EXEC iter 1 score 0 python_exec score=0.00 3.5s

DECIDE iter 1 refine 3.5s

GENERATOR iter 2 1215 chars 5.37s

PYTHON_EXEC iter 2 score 1 python_exec score=1.00 5.4s

DECIDE iter 2 5.4s

LIVE METRICS

Iterations 2

LLM calls 3

Tokens 2,830

Elapsed 5.4s

Agents generator

Best score 1

Validation for code runs the program in a sandboxed subprocess — it is not an LLM opinion. Every stop is bounded by an explicit budget.

5.2 The most important line in this report

Look at the two validator lines above. **The executed tests said FAIL. The LLM critic said PASS.** The LLM was wrong and the program execution was right.

That is Huang et al.'s result reproduced live on our own system, and it is the entire justification for building tool-based validation instead of asking a model to review itself. It is also RQ3 made visible: the interface flags it as a **FALSE ACCEPT**.

5.3 The validator caught a defect in our own benchmark

On the first sweep one task scored **0.75 under every condition**. Identical partial scores across all five is itself a signal — a model failure would vary between conditions, a task failure would not.

The generated code was **correct**. The task asked for a run-length encoding *only when strictly shorter*; for `'aaabbc'` the encoding `'a3b2c1'` is six characters against six, so returning the original was right. Our own test demanded the encoding. **The benchmark was wrong, not the model.** Task rewritten, affected runs deleted, all conditions re-run.

Graded scoring made it visible; the trace made it diagnosable. A project whose entire output is a comparison table must be able to tell *"the system is wrong"* from *"the measurement is wrong"*.

5.4 Measured results so far

25 task-runs with the real API, all five conditions, computed from `traces/smoke.jsonl` :

Condition	n	Pass rate	Mean tokens	Mean calls
A — single agent	5	100%	240	1.0
A+ — budget-matched single	5	100%	431	1.6
B — generator ↔ validator	5	100%	240	1.0
C — fixed pipeline	5	100%	240	1.0
D — adaptive	5	100%	561	2.0

Read this honestly. Every condition solves every task, so these numbers **cannot separate the conditions on correctness**. The smoke set is a development fixture, not a benchmark — the tasks are easy enough for one model call, which is what a smoke set is for. The correct conclusion is that the measurement pipeline works end to end.

What the cost column already shows is the shape of the problem: at identical correctness, **the adaptive condition spent 2.3× the tokens of the single agent**, because its planner adds a call before any work begins. If that gap does not buy correctness on harder tasks, it is cost with no return — exactly the deflationary result this project exists to test rather than assume.

5.5 Engineering status

48 automated tests pass offline, with no API key. They cover budget stops, cycle detection, sandbox timeout and environment isolation, graded partial credit, cache-key sensitivity, test-designer rejections, the budget-matched baseline never using the validator to choose, the metrics module refusing to aggregate mock runs, the demo-token gate, and a full streaming run over HTTP.

6. What comes next

Phase	Dates	Work
Phase 2	11 – 31 Oct	Full benchmark datasets (HumanEval+/MBPP+, Spider, GSM8K); SQL family and its execution validator; evaluate condition D properly; cross-model reviewer experiment
Phase 3	1 – 10 Nov	Frozen benchmark sweep; comparative results; ablation on iteration count (RQ4) and validator type (RQ3); failure classification using MAST; report
Final	23 – 27 Nov	Report, demonstration, viva

6.1 What the datasets are and why those three

Dataset	What it is	Verifier strength
HumanEval+ / MBPP+ (EvalPlus)	164 + ~400 Python problems. EvalPlus adds ~80× more tests than the originals, which were weak enough to let wrong code pass	Strong — per-test partial credit
Spider	~200 databases, English → SQL, checked by running the query and comparing results	Medium — final result only
GSM8K	Grade-school maths word problems, checked on the final answer	Weak — no usable intermediate signal

Three families, not one, because **that axis is the experiment**. One family gives an anecdote; three give a line. This is what answers RQ2.

6.2 A realistic extension, if scope grows

Not application generation. The honest next step is **multi-file Python projects** — a small CLI or package where the system writes a pytest suite for the whole package. That stays verifiable, so the research claim survives. This is the direction SWE-bench points in. It is out of scope for this semester.

7. Limitations, stated plainly

- Only function-level tasks are verifiable. Applications, UIs and servers are declined by design (§2.1).
- Model-written tests are not ground truth. Used for live tasks only; never mixed into benchmark results.

- **Majority voting is weaker for code than for maths.** Two correct programs rarely match character for character, so A+ is a weaker opponent on code. This is reported as a limitation and is itself part of the result.
- **Windows lacks POSIX resource limits**, so the sandbox there enforces only a timeout. Final experiment runs should be done under WSL or Docker.
- **Benchmark contamination** is a known risk for HumanEval and GSM8K; EvalPlus's extra tests reduce but do not remove it. Declared in threats to validity.
- **One model family (Gemini)**, so results may not transfer — which is exactly why the cross-model reviewer experiment is in Phase 2.
- **One API key currently provisioned.** The first sweep already retired it once on a rate limit, which the router handled by design.

Appendix A — How to test the system

Open the demo link, type a task, press **Run**. For the live demo use **condition C**: it puts three distinct agents on screen (generator → critic → repair), which is the shape the project review asked for.

Tasks verified on the deployed instance

Task	Family	Result observed
Evaluate an arithmetic expression with <code>+</code> <code>-</code> <code>*</code> <code>/</code> , correct precedence, no parentheses, no <code>eval</code>	code	0.80 → refine → 1.00 (best demo; under C the critic false-accepted)
Fully justify a list of words into lines of exact width	code	0.80 → refine → 1.00
Convert a Roman numeral to an integer, handling <code>IV</code> , <code>IX</code> , <code>XL</code> , <code>CM</code>	code	1.00 first attempt
Merge overlapping intervals	code	1.00 first attempt
n-th Fibonacci number	code	1.00 first attempt
Second largest distinct value in a list	code	1.00 first attempt
Palindrome ignoring case and punctuation	code	1.00 first attempt
7 pens at ₹12, paid ₹100 — change? (expected answer <code>16</code>)	math	1.00, accepted
"Make a sudoku game as a Flask app"	code	Declined , with the reason

A first-attempt pass is correct but undramatic. **To show the loop working, use one of the first two.**

What to watch, in order

Event	What it proves
<code>TEST_DESIGNER</code> – entry=solve, 5 tests	The task was unseen; tests were written for it on the spot
<code>GENERATOR</code> iter 1	First attempt
<code>PYTHON_EXEC</code> score=0.80	Code actually executed — 4 of 5 tests passed
<code>LLM_CRITIC</code> passed=True	The model's opinion, recorded but never authoritative
<code>DECIDE</code> refine	The policy decided, not the model
<code>REPAIR</code>	A different agent re-solves
<code>PYTHON_EXEC</code> score=1.00 → accept	Fixed

Other things worth demonstrating

- **The budget really stops it.** Ask for something impossible ("a function that solves the halting problem"). It stops on `max_iterations` or `no_improvement` instead of looping — and says which.
- **Maths is the weak verifier.** Run a maths task with its expected answer. The validator can only say right or wrong, never which step was wrong. That contrast with `python_exec` is RQ2 in one screen.
- **It declines what it cannot check.** Ask for an app. The refusal is the point.
- **Offline demo** runs the whole loop with a scripted provider and no API key; runs made this way are tagged `mock` and the metrics module refuses to aggregate them into results.
- **Replay** re-streams a stored real run with original timings, clearly labelled, so the demonstration survives a dead key or venue wifi.

Local

```
pip install -r requirements-dev.txt
pytest -q                                # 48 tests, no API key needed
python -m uvicorn arbiter.web.app:app --reload # http://127.0.0.1:8000

python -m arbiter.bench.runner --dataset smoke --conditions A B C
python -m arbiter.bench.runner --dataset smoke --conditions D --adaptive
python -m arbiter.bench.runner --dataset smoke --conditions A+      # run D first
python -m arbiter.bench.metrics traces/smoke.jsonl --by-family
```

Appendix B — Work distribution

Member	Owns
Gautam Lasgotra (23BCS032)	System architecture, orchestration loop and decision policy, budget governor, provider abstraction with caching and key rotation, sandboxed execution, web interface and streaming, deployment, integration
Chirag Attri (23BCS024)	Agent role library and prompt design, structured output schemas, test-designer agent, LLM critic, provider integrations
Aniket Kundal (23BCS015)	Benchmark task format and datasets, execution validators, metrics module, resumable experiment runner, test suite, failure classification